

Conflicting Guidance on Chunk Sizing (Demo)

Position A: Small Chunks

One school of thought in this demo knowledge base holds that small chunks, on the order of 200 tokens, produce the best retrieval accuracy for question answering, because each chunk stays tightly focused on a single claim and the embedding is not diluted by unrelated surrounding text.

Position B: Large Chunks

A second, contradicting position in this demo knowledge base holds that large chunks, on the order of 1500 tokens, are preferable, because they preserve more surrounding context and reduce the chance that a claim is split awkwardly across a chunk boundary, even though this can dilute the embedding.

These two documents deliberately disagree with each other so that the evaluation dataset can test whether the system surfaces the contradiction explicitly rather than silently picking one side.